



Emilio Quezada

Junior AI Engineer

Contact

+49 177 705 6948

equezadanavarro@gmail.com

Berlin, DE

About Me

Junior AI Engineer specializing in RAG and agentic systems. Six years of operational reliability behind every system I build.

Skills

- **AI & LLM Core:** Generative AI, LLMs, Retrieval-Augmented Generation (RAG), Agentic AI / Multi-Agent Systems, Prompt Engineering, LLM Evaluation, Text Preprocessing & Embeddings
- **AI Frameworks & Orchestration:** LangChain, Hand-rolled Agent Orchestration (multi-hop tool-calling on raw SDKs), Model-agnostic API Integration (OpenAI, Anthropic, Google Gemini), n8n
- **Evaluation & Observability:** LangSmith, Langfuse, Custom Evaluation Harnesses, Hallucination/Grounding Testing
- **Retrieval & Vector Data:** ChromaDB, Supabase, Embeddings, Semantic search
- **Backend & Data:** Python, Flask, FastAPI, REST API, SQL, PostgreSQL, SQLAlchemy, JavaScript, HTML/CSS, Jinja2
- **Cloud & Infrastructure:** AWS (Lambda, RDS, S3, CloudWatch, IAM), Terraform, AWS CDK, Docker, GitHub Actions CI/CD, OIDC Keyless Authentication
- **Security:** JWT auth, Row-level security, Encrypted-field Search / Blind-Indexing
- **Practices:** TDD, Version Control (GitHub), Incident Response & Root-Cause Analysis, Technical Documentation
- **AI-Native Dev Tools:** Claude Code, Cursor, Codex

Experience

AI Engineer Intern

Jul 2026 - Present

Kudan - Tokyo JP

- Architected a hand-rolled multi-hop tool-calling agent loop against the raw Anthropic SDK, enforcing per-request JWT-scoped row-level security and encrypted-field search via a custom blind-index scheme - now the foundation of a conversational AI feature being integrated and tested inside Kudan's PRISM Cloud, a spatial data platform with an established enterprise customer base.
- Built and evaluated 11 production LLM tools for a spatial data conversational agent, achieving 100% grounding and zero hallucinations across a 16-case automated evaluation harness, with full LLM/tool-call observability via self-hosted Langfuse tracing.
- Deployed a production AI service to AWS (Terraform, Docker Compose, GitHub Actions CI/CD with keyless OIDC authentication), root-causing a live outage to a misconfigured logging subsystem via CloudTrail analysis and shipping a permanent fix that cut disk usage from 33GB to 561MB and restored automated deployments.

Operations Manager

Oct 2018 - Jul 2025

Pirate.com - Berlin DE

- Led a cross-functional team of 3 Site Managers spanning IT, Finance, and HR to align global strategic goals, cutting regional rollout implementation time by 50%.
- Developed and executed QA audit protocols and delivered operational training as subject-matter expert, reducing contractor fault-ticket rates by 20% and standardizing service protocols across global teams.
- Applied troubleshooting and diagnostic protocols across multi-site operations while managing operational budgets, cutting technical downtime by an average of 2 hours/week over six consecutive months.

Loc Audio Programmer

Jan 2016 - Dec 2017

Electronic Arts - Madrid ES

- Standardized technical integration and configuration of thousands of localized audio assets across 15 languages for four AAA titles: FIFA 17, Battlefield 1, FIFA 18, and UFC 3, ensuring correct sample-to-context mapping across data banks.
- Executed rigorous QA and batch validation processes for high-volume localized audio data, maintaining consistent quality standards while managing testing, bug fixing, and asset tracking through Perforce version control.
- Analyzed multi-language production pipelines to proactively flag technical risks and bottlenecks ahead of release, informing project-planning decisions across concurrent AAA productions.

Projects

Franz 🤖 - RAG Studio Assistant

Aug 2026 - Sep 2026

[GitHub Repository](#)

- Built a guard → retrieve → generate RAG pipeline (LangChain, Chroma, Claude) with a fully headless backend, hitting 1.00 recall@5 / 0.98 MRR across a 21-case retrieval evaluation.
- Built a regex-first, LLM-fallback topic guard that cuts unnecessary API calls and auto-captures leads into Supabase.
- Instrumented the pipeline with LangSmith tracing and a custom eval suite covering retrieval and answer quality.

Resonate 🌿 - AI Journaling App

Dec 2025 - Mar 2026

[GitHub Repository](#)

- Architected a 5-agent LangChain system: entry analysis, safety screening, non-repeating music curation, and chat, orchestrating OpenAI, Gemini, and Tavily behind one Flask backend.
- Built a fail-closed safety pipeline that detects self-harm and hate-speech signals in real time, routing flagged entries to crisis resources instead of storing them.
- Deployed a two-AZ AWS architecture via Terraform. EC2 running Docker/Gunicorn in a public subnet, RDS PostgreSQL locked to a private subnet with no public database access [aws-deployment branch]

MovieWeb App 🎬 - Multi-user web application

Nov 2025

[GitHub Repository](#)

- Built a full-CRUD, multi-user Flask app with a one-to-many SQLAlchemy ORM model, isolating each user's movie list and integrating the OMDb API for real-time metadata.
- Implemented custom error handling and environment-based API key management across 6 REST routes, following a separation-of-concerns architecture.

Education & Certifications

- Masterschool (MSIT, AZAV-certified): Software Engineering, specializing in AI & Cloud Engineering, Jul 2025 - Sep 2026
- AWS Certified Cloud Practitioner: Amazon Web Services, Jun 2026
- Agentic AI Certification: DeepLearning.AI, Feb 2026
- Retrieval Augmented Generation (RAG) Certification: DeepLearning.AI, Feb 2026
- Audio Engineering: SAE Institute, Sep 2015 - Sep 2016
- Bachelor's Degree, Music Performance and Composition: Conservatori Liceu, Sep 2008 - Sep 2013

Language

• Spanish (Native) | English (C2) | German (B1) | Portuguese (B1)